

# Architecture Decision Record: BigCorp Integration

ADR-001 Status: Approved Date: 2024-03-15

## Context

BigCorp integration requires processing large volumes of data (500k records per import) with strict security and performance requirements. We need to design a scalable, secure, and maintainable architecture.

## Decision

We will implement an event-driven architecture using AWS services with the following components:

### API Layer

- API Gateway for request handling
- Custom authorizer for HMAC validation
- Rate limiting via Redis cluster

### Processing Layer

- SQS for job queue management
- Lambda functions for async processing
- DynamoDB for job status tracking

### Real-time Updates

- WebSocket API for progress updates
- Redis pub/sub for worker communication

### Security Layer

- HMAC authentication
- AES-256 encryption at rest
- Customer-specific encryption keys

# Alternatives Considered

## Synchronous Processing

✗ Rejected due to: - Timeout limitations - Resource inefficiency - Poor user experience for large imports

## Self-hosted Queue

✗ Rejected due to: - Operational overhead - Scaling complexity - Reliability concerns

## Direct Database Import

✗ Rejected due to: - Lack of validation - Security concerns - No progress tracking

# Consequences

## Positive

- Scalable to handle varying load
- Built-in fault tolerance
- Real-time progress updates
- Clear separation of concerns
- Easy to monitor and debug

## Negative

- More complex infrastructure
- Multiple components to maintain
- Higher AWS costs
- Need for careful error handling

# Implementation Details

## Request Flow

```
Client -> API Gateway
        -> HMAC Validation
        -> Rate Limit Check
        -> Create Job (DynamoDB)
        -> Queue Message (SQS)
        -> Return Job ID
```

## **Processing Flow**

SQS -> Lambda Worker  
-> Process Batch  
-> Update Progress (DynamoDB)  
-> Notify WebSocket  
-> Handle Errors

## **Monitoring Points**

1. API Gateway metrics
2. Lambda execution stats
3. Queue depth
4. Processing speed
5. Error rates
6. WebSocket connections

## **Security Considerations**

### **Authentication**

- HMAC signatures required
- API keys with expiration
- IP whitelisting option

### **Encryption**

- TLS 1.3 in transit
- AES-256 at rest
- Key rotation policy

### **Audit**

- CloudWatch logs
- CloudTrail for API calls
- Custom audit events

## **Performance Targets**

### **Throughput**

- 1000 req/min API rate
- 5 concurrent imports
- 100k records/10min

## **Latency**

- API response: < 100ms
- WebSocket updates: < 1s
- Job status: < 500ms

## **Cost Considerations**

- Lambda execution time
- API Gateway requests
- WebSocket connections
- DynamoDB throughput
- SQS message volume

## **Future Considerations**

1. Multi-region deployment
2. Customer-specific queues
3. Enhanced monitoring
4. Automated scaling
5. Backup strategy

## **Team Impact**

- New AWS services to learn
- WebSocket implementation
- Security best practices
- Performance optimization

## **Dependencies**

1. AWS Account
2. Redis Enterprise
3. Security review
4. Performance testing

## **Success Metrics**

1. Processing speed
2. Error rates
3. Customer satisfaction
4. System stability
5. Operational overhead

# Rollout Plan

1. Infrastructure setup
2. Security review
3. Load testing
4. Gradual customer onboarding